

GAMBIT

Stockfish for deals. The negotiation cockpit that thinks like your opponent — and lets you watch it think, or rewire how it thinks.

0. The thesis, in one paragraph

Every transactional negotiation is a game of incomplete information against a counterparty whose moves are *not actually random* — they're a function of how that counterparty has negotiated in the past. Senior partners build this intuition over twenty years of deals. Juniors, AEs, founders, and in-house GCs don't have it, and they bleed money and time because of it. Gambit externalizes that intuition: a *counterparty model* (the Ghost) built from the other side's past contracts, plugged into a chess-engine-style move evaluator (the Board), wrapped inside a cockpit that handles compliance, drafting, sign-off, and execution in one continuous loop. **And then we go further.** Two toggles let you crack the engine open: a **War Room** view that visualizes the multi-agent debate that produced every move, and an **Architect** view — a drag-and-drop workflow canvas where firms rewire the engine into their own house playbook. The redline lands in Gmail. Two minutes later, the contract is countered, compliance-checked, CEO-approved, signed, invoiced in Stripe, and confirmed in Slack — in English or French. *We didn't build a contract assistant. We built an opponent, the war room around it, and the OS underneath it.*

1. The product, in three modes

This is the structural innovation. Gambit isn't one screen. It's one product with three views, toggled by a tab strip across the top.

Mode	Audience	What you see	When to demo
Cockpit	Closers (Sarah, in-house GCs, AEs, founders)	The chess board. Eval bar, game tree, ghost, walkaway line, side panels for compliance and execution. The streamlined hero experience.	Default. ~80% of the demo.

War Room	Lawyers under malpractice pressure, audit / sign-off scenarios	The multi-agent debate visualized as a council. Agents argue, vote, influence the move. Every recommendation has a paper trail.	Toggle 1. ~10s flash of depth.
Architect	Firm administrators, GCs, ops leads, anyone building a house playbook	A drag-and-drop canvas. Agents as blocks. Edges as data flow. Sliders for risk, urgency, must-win clauses. Save as playbook.	Toggle 2. The platform play.

Why three modes is the right shape: The cockpit alone is a *product*. The War Room makes it *trustworthy* (every move auditable). The Architect makes it a *platform* (every firm rewires it). Most teams in the room will build one of these. We build all three, integrated, and pitch them as a single coherent stack.

The demo is the cockpit. The toggles are the "*oh wait, there's more*" reveal that turns a 7/10 demo into a 10.

2. Why this is the winning shape (vs. everything else we considered)

We had five serious candidates on the table. Each was strong; each was incomplete. The fusion is what wins.

Candidate	Strength	Why it loses alone
ClauseOS (drag-and-drop agent canvas)	Visual, agentic, customizable, sponsor-friendly	Drag-and-drop agent builders are the dominant UI pattern of 2024–26 (n8n, Langflow, Dust, Vellum, Flowise). Judges have seen the canvas fifty times. Alone, it's not novel. Inside Gambit, as Architect mode, it's the platform layer.
Stockfish-eval / chess game tree	Iconic visual, universally legible, viral screenshot	Without an opponent, what is the eval bar evaluating against? It's a math toy without an adversary. Inside Gambit, the Ghost is the opponent.

Counterparty Ghost (built from past contracts)	Genuinely novel — nobody in legaltech does this. Mirrors what senior partners actually do.	Conceptually thin demo without a board to play on. Inside Gambit, the Ghost moves on the chess board.
Clausefish (multi-agent debate)	Sophisticated, agentic, auditable	Multi-agent debates demo as walls of text. Boring on stage. Inside Gambit, it's the toggleable War Room view — visualized as a council, not a chat log.
Sarah / WhatsApp / voice wrapper	Humanizes the product. Real user, real pain.	A texting demo isn't theatrical. Inside Gambit, Sarah is the protagonist of the cockpit demo.

Each is a 7/10 alone. Fused, with the new layers (TrueSight, Law 25, Execution Engine, Supervisor PDF, French, Calendar, Voice approval), it's a 10. Nothing got cut. Everything became a layer in a coherent product.

3. The protagonist: Sarah

The pitch opens on her, not the product. Every law student in that room has been the lawyer Sarah is begging. Every business judge has been Sarah.

4:00 PM. Last day of Q4. Sarah is an AE at a Series B SaaS company. She needs one signature — a SaaS MSA from MegaCorp Cloud Services — to hit her number. MegaCorp's legal team just emailed back twelve redlines. Sarah's own counsel is swamped and tells her they need 48 hours. Her quarter ends in two.

Pause. *"This is the moment Gambit was built for."*

The product now lands on a person, not a contract.

4. Mode 1: The Cockpit — Sarah's view (the hero demo)

This is the streamlined experience. End-to-end, in the order Sarah encounters it.

4.1 The trigger — Gmail and Slack listening

Gambit watches Sarah's inbox (Gmail OAuth) and her Slack channels. The moment a redlined contract arrives — recognized by attachment type, subject line patterns, and content

classification — the cockpit lights up. No upload step. The deal walks in the door and Gambit is already at the table.

4.2 The pile — what they sent us

The redlined MSA renders in the center as a Word-doc pane: twelve highlighted changes, struck-through text, marginal comments. On the right edge, a chess **eval bar** materializes and tilts to **-2.4 in MegaCorp's favor**. The gut-punch moment: *we're losing this deal and we don't even know it yet*.

4.3 The deadline — Calendar context made physical

A small pane in the corner pulls from Google Calendar. Sarah has a **board meeting at 6:00 PM**. Gambit flags: *"Deal must close in 1h 47m to clear before your board meeting."* The judges feel the deadline in their gut. (Calendar integration: real, with conflict detection.)

4.4 The Ghost — the moment the demo wins

Sarah clicks **"Build Opponent."** Gambit ingests three of MegaCorp's past public MSAs (we pre-load these as fictional but defensibly-sourced precedents). A counterparty card materializes, animating in like a fighting-game roster screen:

MEGACORP LEGAL

ELO: 2,180 · Hardline

Playstyle: Aggressive on IP. Conservative on liability.

Fights to the end on: IP indemnity · Audit rights · Data residency

Often concedes: Payment terms · MAC clauses · Notice periods

Walks the table when: Liability cap < \$3M

Trained on 3 prior deals · 47 clauses analyzed

This is **the screenshot**. This is **the meme**. A character forming on screen — like a fighting-game opponent profile — for a *legal counterparty*. Nobody has done this in legaltech. Non-technical judges feel it instantly. Technical judges recognize it as RAG + structured extraction over precedent contracts.

4.5 The bloom — the chess board for deals

The current position renders as a **decision tree**, blooming outward from the root. Three candidate counter-moves branch first, each tagged with chess engine notation:

- **!! Brilliant** — *"Concede Net 60, hold liability cap at \$5M, tighten audit rights."* → 82% close probability, +\$280K retained value.

- **?! Risky** — *"Push back on all twelve redlines."* → 31% close probability, but +\$420K if it lands.
- **?? Blunder** — *"Accept their full markup."* → 96% close probability, but you eat \$2M of latent risk.

Hover a branch and **MegaCorp's likely response materializes** — generated live by the Ghost. Branch deeper, you see counter-counters two ply down. The eval bar **tracks live as Sarah hovers** different lines.

4.6 The walkaway line — the legal-defensibility moment

A horizontal red line draws across the bottom of the tree. *"Below this line, MegaCorp walks."* Three citations pop: actual quoted clauses from past MegaCorp MSAs. *"In the Acme deal (2023), MegaCorp walked when the cap dropped below 2.1× ARR. Here. And here. And here."*

This is where Gambit goes from "cool toy" to "thing I might actually use." Every prediction is footnoted to evidence.

4.7 TrueSight — the anti-hallucination firewall

As the engine drafts a counter-redline, one suggestion references *"per the holding in Beaulieu v. Provincial Insurance, 2021 QCCA 412."* The text appears, then **glows red and strikes through itself**. A badge replaces it: *"⚠ Hallucination caught. Case not found in CanLII. Substituted with verified citation: Bhasin v. Hrynew, 2014 SCC 71."*

Lawyers in 2026 are *terrified* of AI — between *Mata v. Avianca*, the Illinois sanction, the *NCR v. KKB* warnings out of Alberta — every law student knows three lawyers who've been burned. A self-auditing AI that catches its own hallucinations against CanLII is the single most credibility-building moment of the demo.

4.8 Law 25 Trapdoor — the Quebec moment

One redline Sarah is considering accepting includes *"tracking technologies, IP-based geolocation, and cross-border data transfers to U.S.-based subprocessors."* The cockpit flashes: **"⚠ LAW 25 TRIGGERED."**

A side panel slides open. Inside: a **drafted Privacy Impact Assessment**, auto-generated from the contract parameters — categories of data, retention, transfer destinations, lawful basis. **In French and English** (toggle at the top — Cloud Translation API, real). The PIA looks like real Quebec compliance work product. It's the moment local judges realize Gambit understands their world.

4.9 Play Best Line — the work product

Sarah clicks **"Play Best Line."** The tree highlights a four-move sequence. The eval bar slides from -2.4 to **+0.8**. A real **.docx** renders in the side pane — track changes, proposed counter-redlines on the actual MegaCorp MSA. Not a summary. Not a chat answer. A sendable, downloadable contract.

4.10 The Supervisor — CEO sign-off layer

Before anything goes out, Gambit assembles a **Supervisor Report** — a clean 2-page PDF designed for the CEO/CFO/GC. Recommended counter, dollar impact, compliance flags, decision tree summary, signature block.

A toggle at the top: **"Legal Mode / Plain-English Mode."** Flip it, and the same document re-renders in business speak — no Latin, no "notwithstanding the foregoing," just *"Here's what they asked for, here's what we're giving back, here's why, here's what could go wrong."*

This makes Gambit usable by *organizations*, not just lawyers. The lawyer drafts the position. The CEO signs it. They speak different languages. Gambit translates between them with a switch.

4.11 Sarah signs — voice in the loop

Sarah's voice memo plays (pre-recorded for demo): *"Looks good. Send it."* The Supervisor report is signed.

4.12 The Execution Engine — contracts as code

This is the closer beat nobody else in the room will build. The recommended terms include *"Net 30 payment of \$9,500/month for 12 months, software tier: Enterprise, seats: 50."* Gambit fires three webhooks:

1. **Stripe** — pending invoice, \$9,500/month Net 30, customer = MegaCorp.
2. **Provisioning** — 50 Enterprise seats pre-allocated in a mock SaaS dashboard.
3. **Notary queue** — for Quebec, contract routed for witnessing where required.

A timeline animates across the bottom: *Email received → Counter sent → Invoice generated → Seats provisioned → Notary queued. Five green checkmarks in real time.*

This is the closing line: **"Most contract AI ends at the redline. Gambit ends at the invoice."**

4.13 Outbound — Slack and Gmail

Slack notification fires in the team channel: *"Gambit recommended counter on MegaCorp MSA. CEO approved. Sent at 5:47 PM."* The Gmail draft is open and pre-filled — *To: MegaCorp Legal, Subject: RE: MSA — Counter-redlines attached*, body pre-written, Word doc attached. Sarah hits Send. The deal is in motion.

Cut to her calendar: **5:47 PM. Board meeting at 6:00 PM. She made it.**

5. Mode 2: The War Room — the multi-agent debate, visualized

This is where you toggle to in the demo when you want to answer the question every law student is going to ask: *"How do I know I can trust the AI's recommendation?"*

5.1 What the user sees, instead of a text wall

Imagine a council. Five agents arranged around a virtual round table, each with a name, an icon, a role, and an **influence meter** filling in real time. They don't argue in paragraphs — they argue in **short speech bubbles** that animate in and fade out, like a Pixar movie.

Agent	Role	Argues for
Counsel	Legal Risk Agent	Indemnity, IP, liability, termination, data privacy. Hates risk.
Closer	Business Goals Agent	Revenue, deal speed, customer relationship, urgency. Hates friction.
Counterpart	Counterparty Simulator	What MegaCorp will accept, push back on, walk on. Powered by the Ghost.
Compliance	TrueSight + Law 25 Watcher	CanLII verification, jurisdictional triggers, regulatory exposure.
The Crown	Judge / Move Scorer	Synthesizes votes. Calls the final move.

5.2 What the debate looks like, visually

For each candidate move (the three branches from the bloom), the council fires:

Counsel: *"Liability cap at \$5M is within tolerance. Vote: ACCEPT."* — speech bubble, 1.2 seconds, fades. **Closer:** *"Net 60 over Net 30 hits Q1 cash flow. Tradeable for cap retention. Vote: ACCEPT."* — fades. **Counterpart:** *"MegaCorp's prior deals show 73% acceptance of Net 60 when cap holds. Vote: LIKELY COUNTER-ACCEPT."* — fades. **Compliance:** *"No Law 25 triggers on this branch. Citations verified against CanLII. Vote: CLEAR."* — fades. **The Crown:** *"Council vote: 4–0 in favor of Move A (Brilliant). Final recommendation locked."* — settles.

Each agent's influence meter fills. A final vote tally appears next to each candidate move. It's **15 seconds, all visual, all auditable.**

5.3 Why this matters legally

The War Room view is *not just a gimmick*. It's a **defense artifact**. Every move Gambit recommends comes with a record of how it was decided — which agents argued for and against, what evidence they cited, how the final vote landed. In a malpractice context, in a sign-off context, in an audit context, this is gold. Lawyers can point at the record and say *"here is why I made this decision."*

In the pitch: ***"Gambit doesn't just give you a move. It gives you the deliberation that produced the move. Every recommendation is auditable. That's not a feature — that's how AI gets used in regulated industries."***

5.4 The optional theatrical flex (Q&A only, if time)

If a judge asks *"can the agents ever disagree?"*, we trigger a **live AI-vs-AI mini-negotiation**. Two instances of the engine, one playing MegaCorp, one playing Acme, negotiate a clause in real time. Eval bar swings. Speech bubbles fly. Judges interject *"what if MegaCorp gets aggressive on IP?"* and watch the bar move. This is the Twitch-stream / chess-broadcast moment if we have time. **We don't put it in the main demo**. It's the flex for Q&A or the wow moment if the demo is going long and we're killing it.

6. Mode 3: The Architect — the drag-and-drop platform layer

This is the second toggle. It's the "but every firm negotiates differently" reveal.

6.1 The canvas

A drag-and-drop workspace, similar in feel to n8n / Langflow / Figma, but built specifically for legal negotiation workflows. The default canvas, when you open Architect mode, **already shows Gambit's stock workflow** rendered as a graph — agents as nodes, data flows as edges, integrations on the periphery.

The user can:

- **Drag new agent blocks** from a sidebar (Spellbook Issue Detector, Counterparty Ghost, Counsel, Closer, Counterpart, Compliance, The Crown, TrueSight, Law 25 Watcher, Translation Layer, Slack Notifier, Gmail Drafter, Stripe Webhook, Provisioning Webhook, Notary Router, Supervisor PDF Generator, and so on).

- **Tune each agent's parameters** with sliders — risk tolerance, urgency, deal size threshold, must-win clauses, acceptable concessions, tone of negotiation, close-probability target.
- **Reconnect edges** to change how data flows.
- **Save as Playbook** — name it *"Stikeman Aggressive Close"* or *"Fasken Civil Law Default"* or *"Bay Street Tech M&A"* — and now it's the firm's negotiation IP.
- **Share playbooks** across the firm (the Slack/Notion analogy: the firm builds a library).

6.2 The 20-second demo

"Out of the box, Gambit works. But every firm negotiates differently."

[Toggle to Architect. The default workflow renders.]

"This is the default. Now let's say we're configuring this for a Bay Street M&A boutique that does aggressive seller-side deals."

[Drag in a new agent block: "Sector-Specific Risk Profile (M&A)." Reweight the Closer agent's urgency slider from 0.5 to 0.8. Drop in a Notary Router for Quebec compliance.]

"Hit save."

[Click "Save as Playbook." A modal: "Stikeman Tech M&A Playbook v1.0." Save.]

"That playbook is now live. Every deal that flows through this firm uses this engine. Multiply that across 200 lawyers. That's institutional knowledge made executable."

20 seconds. Tight. Punchy. **Gambit is no longer a tool. It's an OS for negotiation.**

6.3 Why this matters strategically

The Architect view is what makes Gambit a **platform**, not a feature. Three implications the pitch needs to hit:

1. **Every firm has a unique negotiation style.** A Stikeman partner negotiates differently than a Norton Rose partner than an in-house GC at a startup. The Architect lets each one configure Gambit to their style.
2. **Playbooks become institutional IP.** A senior partner's mental model — encoded as a playbook — can be shared across juniors. **Tacit knowledge becomes transferable.**
3. **The platform extends.** New compliance regimes (EU AI Act, federal AIDA, sector-specific regs) are new agent blocks. Gambit grows with the regulatory landscape without re-architecting.

In the pitch: ***"The cockpit is the product. The Architect is the platform. The playbook is the IP."***

6.4 The sponsor angle

Spellbook is a clause-detection company. Inside Architect, Spellbook is the *first node* in the default workflow — every Gambit deal flows through Spellbook for issue detection before the engine takes over. We're not replacing Spellbook. We're putting them at the front of the pipeline. **The demo should pause on this node and call it out by name.** Sponsor track wins are won by making the sponsor look indispensable.

7. The bilingual layer — not a feature, a posture

Quebec hackathon. Half the judges will be French-first. The bilingual feature isn't a translation toggle in a settings menu. It's a *posture*: Gambit is built to operate in both languages natively. The PIA generates in French because Quebec compliance work product *is* in French. The Supervisor renders in either. Slack notifications, Gmail drafts, even chess engine notation — all of it switches.

In the pitch: ***"We didn't translate Gambit into French. We built Gambit bilingually. The compliance workflow is in French because Law 25 is in French. The contract is in English because that's what Sarah's counterparty sent. Gambit handles both at once, on the same screen."***

8. The features, mapped to demo beats

So nothing gets lost: every feature is in here, and every feature is a *beat*, not a bullet in an appendix.

Feature	Where it lives	Demo beat	Mode
Gmail integration	Trigger + outbound draft	Beats 1, 13	Cockpit
Slack integration	Outbound notification	Beat 13	Cockpit
Google Calendar (with conflict detection)	Deadline pressure	Beat 3	Cockpit
Counterparty Ghost	Core engine	Beat 4	Cockpit
Chess eval bar + game tree	Core interface	Beats 2, 5, 9	Cockpit

Walkaway line w/ citations	Legal defensibility	Beat 6	Cockpit
TrueSight (anti-hallucination via CanLII)	Compliance layer	Beat 7	Cockpit
Law 25 Trapdoor + auto-PIA generator	Compliance layer	Beat 8	Cockpit
French/English bilingual	i18n posture	Beats 8, 10, throughout	All modes
Word doc output (real .docx)	Work product	Beat 9	Cockpit
Supervisor PDF + plain-English toggle	Sign-off layer	Beat 10	Cockpit
Voice memo approval	Human-in-the-loop	Beat 11	Cockpit
Stripe invoice trigger	Execution engine	Beat 12	Cockpit
Software provisioning webhook	Execution engine	Beat 12	Cockpit
Notary distribution queue	Execution engine (Quebec)	Beat 12	Cockpit
Multi-agent debate (Council/Crown)	Auditability + transparency	Toggle 1	War Room
AI-vs-AI mini-negotiation	Theatrical flex (Q&A only)	Optional	War Room
Drag-and-drop agent canvas	Customization	Toggle 2	Architect
Tunable agent parameters	Per-firm playbook	Toggle 2	Architect
Save as Playbook	Institutional IP	Toggle 2	Architect
Spellbook integration	Issue detection	Implicit + named in Architect	All modes
SSO mock accounts (Sarah, MegaCorp Legal, CEO)	Demo polish	Throughout	All modes

Every feature on every list is here. Nothing was cut.

9. The full demo arc — exact pacing

This is a 3-minute demo (the standard at most hackathons). The cockpit beats are tight — 5–10 seconds each. The toggles are quick depth-reveals. The close lands the pitch.

Time	Beat	Mode
0:00 — 0:15	Sarah cold open. Pain. Stakes. <i>"This is the moment Gambit was built for."</i>	—
0:15 — 0:25	The pile. Twelve redlines. Eval bar tilts to -2.4.	Cockpit
0:25 — 0:35	Calendar context. Board meeting at 6 PM. Deadline made physical.	Cockpit
0:35 — 0:55	The Ghost forms. Counterparty card materializes. The screenshot moment.	Cockpit
0:55 — 1:25	The bloom. Three branches with chess notation. Hover responses. Eval bar tracks live.	Cockpit
1:25 — 1:40	Walkaway line + cited evidence. <i>Three deals where they walked. Here. Here. Here.</i>	Cockpit
1:40 — 1:55	TrueSight intercepts a hallucinated citation. Glowing red strikethrough. Substituted with verified case.	Cockpit
1:55 — 2:15	Law 25 triggers. PIA auto-generates in French + English. <i>The Quebec moment.</i>	Cockpit
2:15 — 2:30	Toggle 1: War Room. Council debates. Speech bubbles. Crown calls 4–0 for Move A. <i>"Every recommendation is auditable."</i>	War Room
2:30 — 2:45	Toggle 2: Architect. Drag a new agent. Tune a slider. <i>"Save as Stikeman Tech M&A Playbook."</i> <i>"This is institutional IP."</i>	Architect
2:45 — 2:55	Back to Cockpit. Play Best Line. Tree highlights. Eval bar to +0.8. Word doc renders.	Cockpit
2:55 — 3:05	Supervisor PDF. Plain-English toggle. Voice memo: <i>"Send it."</i>	Cockpit
3:05 — 3:20	Execution Engine. Stripe invoice. Seats provisioned. Notary queued. Slack pings. Gmail sent.	Cockpit
3:20 — 3:30	Cut to calendar. 5:47 PM. She made the board meeting. Close.	—

Three minutes, thirty seconds. Every feature has a beat. Every beat has a visual. Nothing dwells.

10. Differentiation, baldly

Roughly 30% of teams will build a redline-explainer chatbot — *which is literally Spellbook's existing product*, so they're dead on arrival. ~20% will build a clause-risk heatmap. ~15% will build a market-standard comparator. ~10% a Q&A chatbot. ~10% an auto-redline generator. ~5% a deal-desk dashboard. ~10% will attempt something genuinely novel and most will be confused executions.

Five of those six categories end at the same demo: an annotated PDF with colored boxes. They will blur together by the fourth pitch.

Gambit is in a different category entirely. The Ghost, the eval bar, the Council debate, the Architect canvas, the live execution chain — none have analogues in the room. **Judges will remember us because our demo *moves* in a fundamentally different shape than everyone else's.**

11. Academic anchoring — the references that win law-student rooms

Most teams will quote zero canonical sources. We will quote three, casually, in one sentence each, and the law-student judges will sit up.

- **Fisher & Ury, *Getting to Yes* (1981)** — BATNA, ZOPA. The eval bar represents Sarah's position relative to her BATNA. The walkaway line is MegaCorp's reservation point.
- **Howard Raiffa, *The Art and Science of Negotiation* (1982)** — decision trees under uncertainty. The bloom *is* a Raiffa tree. Each branch is probability-weighted. Each leaf has utility.
- **Roger Fisher's principled negotiation framework** — informs the agent priorities (interests over positions). The Council's structure is a principled-negotiation cabinet.

The single sentence in the pitch: ***"Under the hood, the engine is a Raiffa decision tree with reservation-value-anchored leaves and ZOPA-bounded branches. The Ghost models the counterparty's BATNA from prior conduct. We're not making this up — we're operationalizing forty-year-old negotiation theory with 2026 tools."***

12. What's real, what we LARP, where the line is

Genuinely real (we will demo these live):

- Gmail OAuth + Slack webhook + Google Calendar API
- Claude API for Ghost generation, move scoring, redline drafting, agent debate
- Real `.docx` output via `python-docx`
- Real PIA generation, translated FR/EN via Cloud Translation
- TrueSight live citation lookup against CanLII
- Stripe invoice creation (test mode, real API)
- Supervisor PDF generation
- React-Flow game tree visualization
- Drag-and-drop Architect canvas (built on React-Flow's existing graph editor — surprisingly fast to ship)
- War Room Council visualization (custom React + Framer Motion)

Theatrical (defensible if asked):

- *"Trained on 10,000 deals" → in reality, 3 hardcoded MegaCorp precedents per counterparty. Defense: "The mechanism is real — RAG over precedent contracts. The corpus is the production challenge. Hackathon corpus is three deals; production version pulls from EDGAR, the firm's deal vault, and counterparty public filings."*
- MegaCorp ELO of 2,180 → vibes-correct composite of stubbornness × win rate × concession depth.
- Voice memo from Sarah → pre-recorded. We say *"voice-enabled approval flow,"* not *"live STT."*
- Some demo branches in the tree are pre-staged so the bloom always lands. Live generation works as fallback in Q&A.
- *"73% accuracy predicting next moves" → never said unless asked. Backtest framing if pressed.*
- SSO mock accounts (Sarah, MegaCorp Legal, CEO) are fake users, but the experience is end-to-end real.

What we cut:

- iMessage. Closed, painful, ten-minute hassle for zero demo upside.
 - Showing the agents' debate as raw text. (Replaced with the Council visualization.)
-

13. The pitch — open and close

Open (15 seconds):

"Every transactional lawyer in this room has stared at a redline and thought 'I have no idea what they're going to come back with.' Senior partners build that intuition over twenty years of deals. We compressed it into ninety seconds. We're Gambit."

Close (20 seconds):

"Spellbook tells you what the redline says. Harvey tells you what the contract means. Gambit tells you what they're going to send next — and the move that wins the deal. And then it sends it, in either official language, with a Privacy Impact Assessment if Law 25 fires, with a CanLII-verified citation, with the CEO's sign-off, with a pending Stripe invoice, and with the Council's full deliberation on the record. Every firm rewires the engine to their playbook. Most contract AI ends at the redline. Gambit ends at the invoice. We didn't build an assistant. We built an opponent — the war room around it — and the OS underneath it."

That last sentence is the one judges write down.

14. Risks and trade-offs

Risk 1: Three modes in 3:30 is ambitious. Mitigation: cockpit beats are tight (5–10s each); toggles are 15–20s flashes, not deep dives. We rehearse mercilessly. Default plan: 3:30 demo. Can compress to 2:45 if needed by cutting Stripe + Notary beats (still keep the timeline animation, just faster).

Risk 2: The Ghost feels gimmicky if the card looks thin. Mitigation: invest in the card. Typography, animation, real numerical detail. This is the screenshot.

Risk 3: The Architect canvas takes time to build. Mitigation: use React-Flow's out-of-the-box graph editor — drag, drop, connect, parameter-tune is mostly free. We pre-load the default workflow so the canvas isn't blank. Hour 7–8 of the build.

Risk 4: Multi-agent visualization could look childish if the avatars are too cute. Mitigation: use professional iconography (think Linear, Vercel, Stripe aesthetic) — clean monochrome icons with colored accent rings, not cartoon faces. The "Council" framing is serious, not cutesy.

Risk 5: A judge asks *"is the Ghost real?"* Mitigation: **it is**. RAG over three pre-loaded contracts with structured extraction. We can run it live on a fourth contract in Q&A. The LARP is the *scale*, not the *mechanism*.

Risk 6: Spellbook panelists feel competed-with. Mitigation: explicit framing throughout. Spellbook is the **first node** in the default Architect workflow. The pitch credits them by name. Make them look indispensable, on purpose.

15. The closer

The single sentence that captures all of this:

"Spellbook reviews the contract. Gambit closes the deal — from inbox to invoice, in two languages, with the Council's record on the file, with the CEO's signature already on it, on the playbook your firm built."

That's the project. That's the pitch. That's the demo.

16. The three locks before we start building

1. **Name lock: Gambit.** Bilingual-friendly (gambit = gambit in French chess), strategic, memorable.
2. **Stack lock:** Next.js + Tailwind + shadcn front-end. **React-Flow** for both the game tree (Cockpit) and the agent canvas (Architect) — same library, double duty. **Framer Motion** for the Council animations. Claude API for engine + Ghost + agents. **python-docx** and **reportlab** for outputs. Cloud Translation API for FR/EN. Slack webhooks + Gmail OAuth + Google Calendar API + Stripe test mode. CanLII for TrueSight. Vercel deploy.
3. **Demo deal lock:** Spellbook's provided fictional MSA + redlines as the live deal. Counterparty = "MegaCorp Cloud Services Inc." (fictional → fabricate three precedent MSAs without legal-name risk). Sarah's company = "Acme Analytics." Three SSO mock accounts pre-staged: Sarah Chen (AE), Marc Tremblay (GC, signs Supervisor), MegaCorp Legal (sender of redlines).

Lock those three and I'll write the **hour-by-hour build plan for the 10-hour sprint**, with explicit fallback gates at hours 4, 6, 8, and 10 — so if anything's behind, we know exactly what to cut without breaking the demo arc.